


► CaVE

**A Cone-Aligned Approach for
Fast Predict-then-Optimize**



Authors



Bo Tang

PhD Candidate
Department of Mechanical &
Industrial Engineering,
University of Toronto



Elias B. Khalil

Assistant Professor
Department of Mechanical &
Industrial Engineering, University
of Toronto

SCALE AI Research Chair
Data-Driven Algorithms for
Modern Supply Chains

Introduction

CaVE (Cone-aligned Vector Estimation) is an efficient and accurate **Decision-focused Learning** approach for **Binary Linear Programs** (BLPs).



End-to-End:

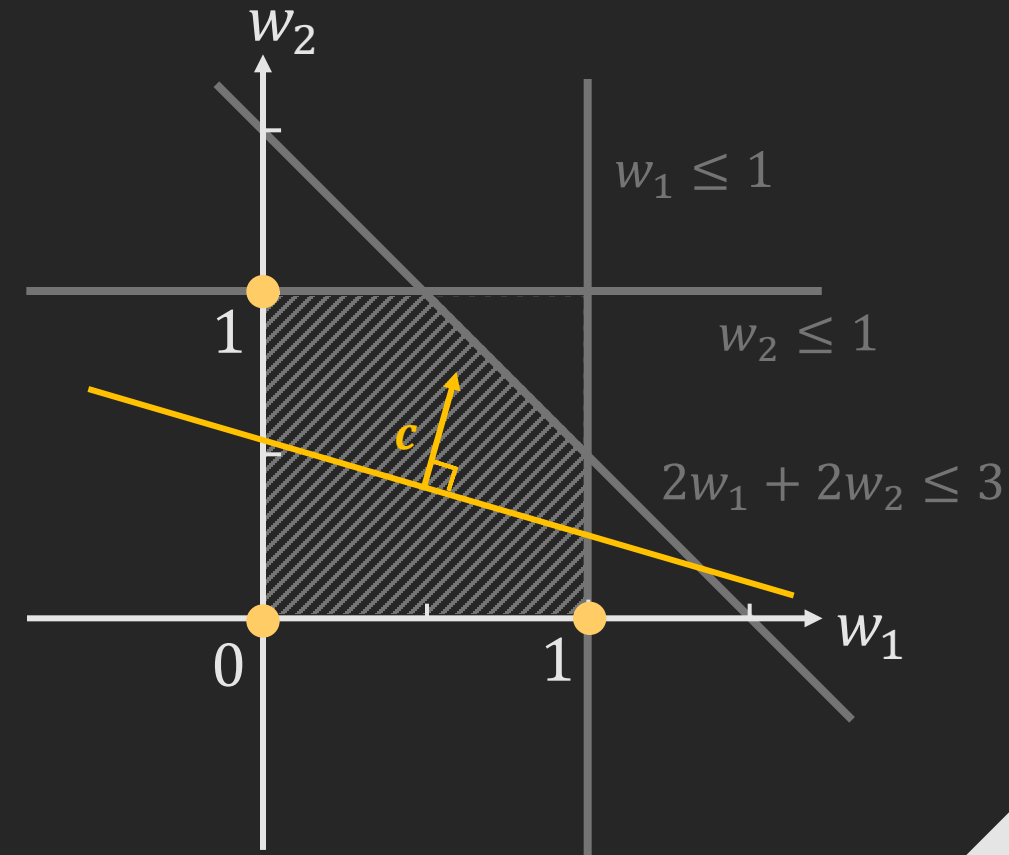
The loss function of CaVE focuses on decision quality.

Efficiency:

The algorithm of CaVE utilizes non-negative least squares (NNLSs) instead of solving BLPs.

Notation

$$\min_{\underbrace{w}_{\text{Decision Variables}}} \left\{ \overbrace{c^T w}^{\text{Linear Objective Function}} : \underbrace{w \in S}_{\text{Feasible Region}}, \underbrace{w \in \{0,1\}}_{\text{Binary Domain}} \right\}$$



Predict, then Optimize

$$\min_w \{ \mathbf{c}_1^\top \mathbf{w} : \mathbf{w} \in S \}$$

$$\min_w \{ \mathbf{c}_2^\top \mathbf{w} : \mathbf{w} \in S \}$$

$$\min_w \{ \mathbf{c}_3^\top \mathbf{w} : \mathbf{w} \in S \}$$

$$\vdots$$

Predict, then Optimize

Unknown
Coefficients

$$\min_w \{ \mathbf{c}_1^\top \mathbf{w} : \mathbf{w} \in S \}$$

$$\min_w \{ \mathbf{c}_2^\top \mathbf{w} : \mathbf{w} \in S \}$$

$$\min_w \{ \mathbf{c}_3^\top \mathbf{w} : \mathbf{w} \in S \}$$

⋮ Identical
Constraints

Predict, then Optimize

Unknown
Coefficients

$$\begin{aligned} \min_w \{ \mathbf{c}_1^\top \mathbf{w} : \mathbf{w} \in S \} \\ \min_w \{ \mathbf{c}_2^\top \mathbf{w} : \mathbf{w} \in S \} \\ \min_w \{ \mathbf{c}_3^\top \mathbf{w} : \mathbf{w} \in S \} \\ \vdots \end{aligned}$$

Identical
Constraints

Observed Feature Vector

\mathbf{x}_1

\mathbf{x}_2

\mathbf{x}_3

\vdots

Predict, then Optimize

Unknown
Coefficients

$$\begin{aligned} \min_w \{ \mathbf{c}_1^\top \mathbf{w} : \mathbf{w} \in S \} \\ \min_w \{ \mathbf{c}_2^\top \mathbf{w} : \mathbf{w} \in S \} \\ \min_w \{ \mathbf{c}_3^\top \mathbf{w} : \mathbf{w} \in S \} \\ \vdots \end{aligned}$$

Identical
Constraints

Observed Feature Vector

\mathbf{x}_1

\mathbf{x}_2

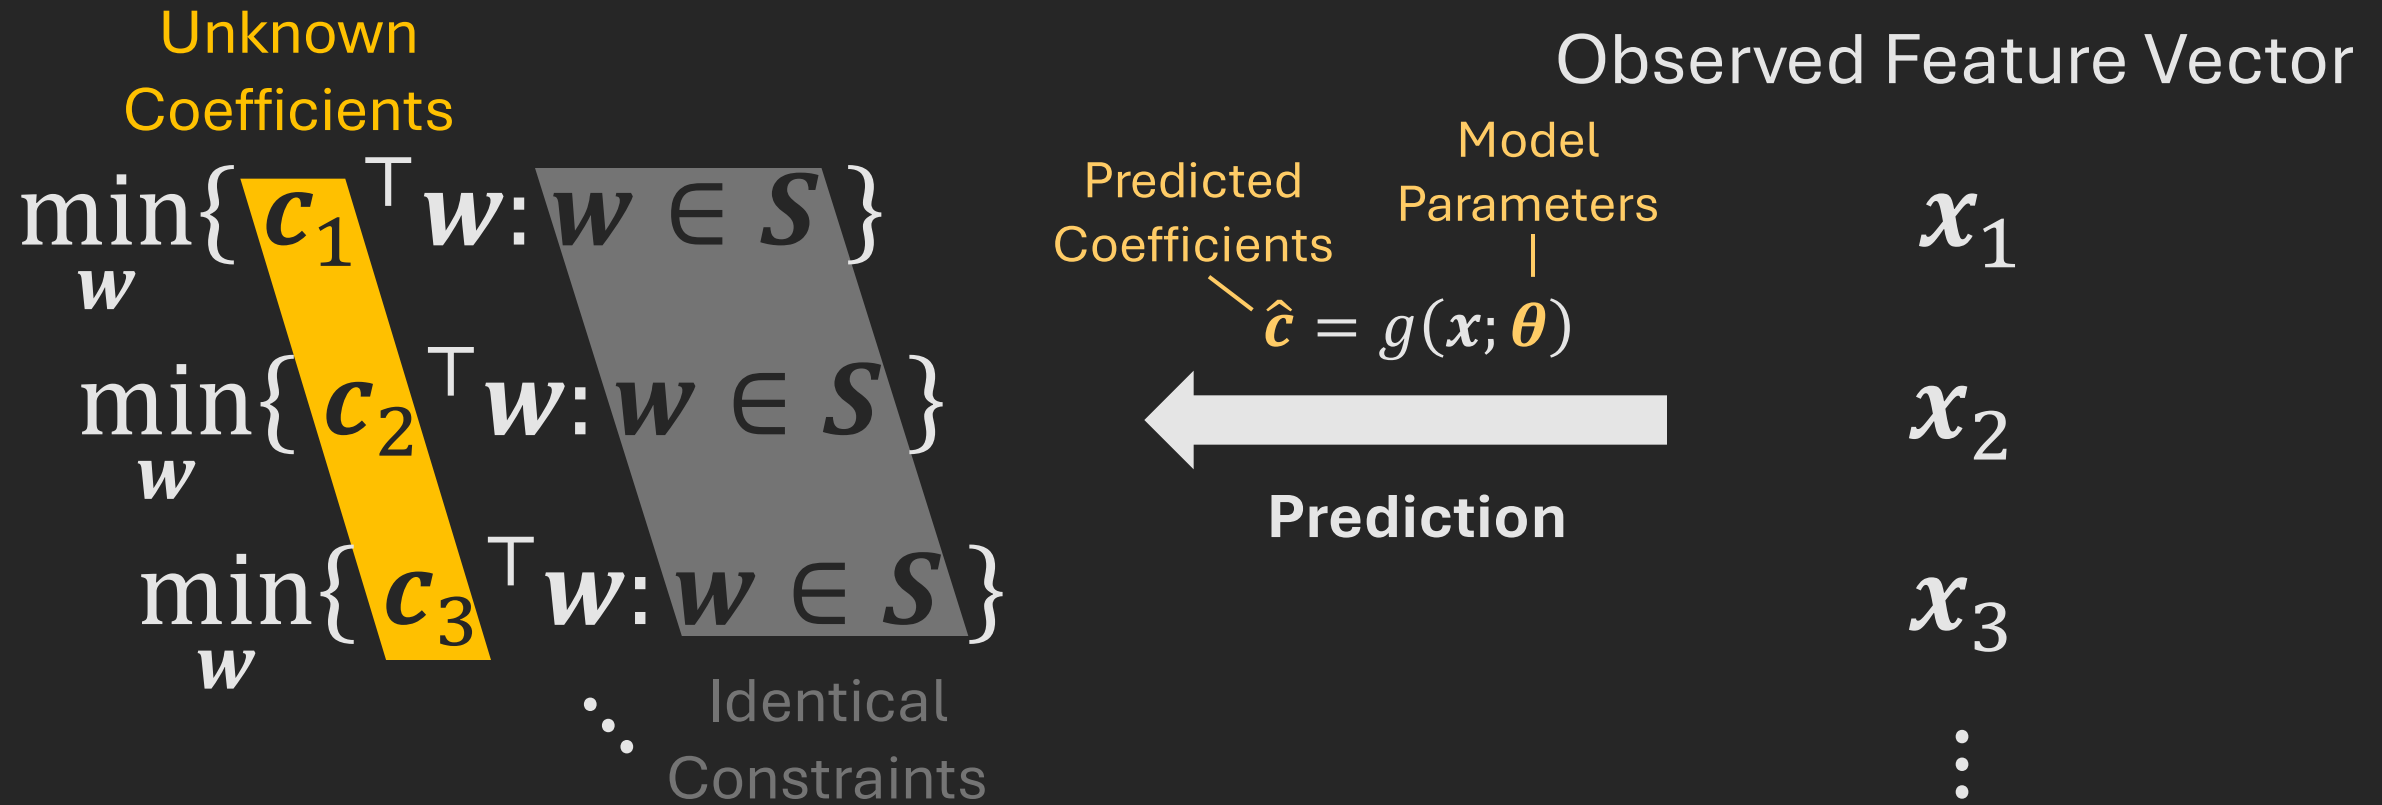
\mathbf{x}_3

\vdots

$$\hat{\mathbf{c}} = g(\mathbf{x}; \theta)$$

Prediction

Predict, then Optimize



Examples



❖ Vehicle Routing



❖ Energy Scheduling



❖ Portfolio Optimization

Examples



❖ Vehicle Routing



❖ Energy Scheduling



❖ Portfolio Optimization

? Unknown Costs: Travel Time, Electricity Prices, Asset Returns

Examples



❖ Vehicle Routing



❖ Energy Scheduling



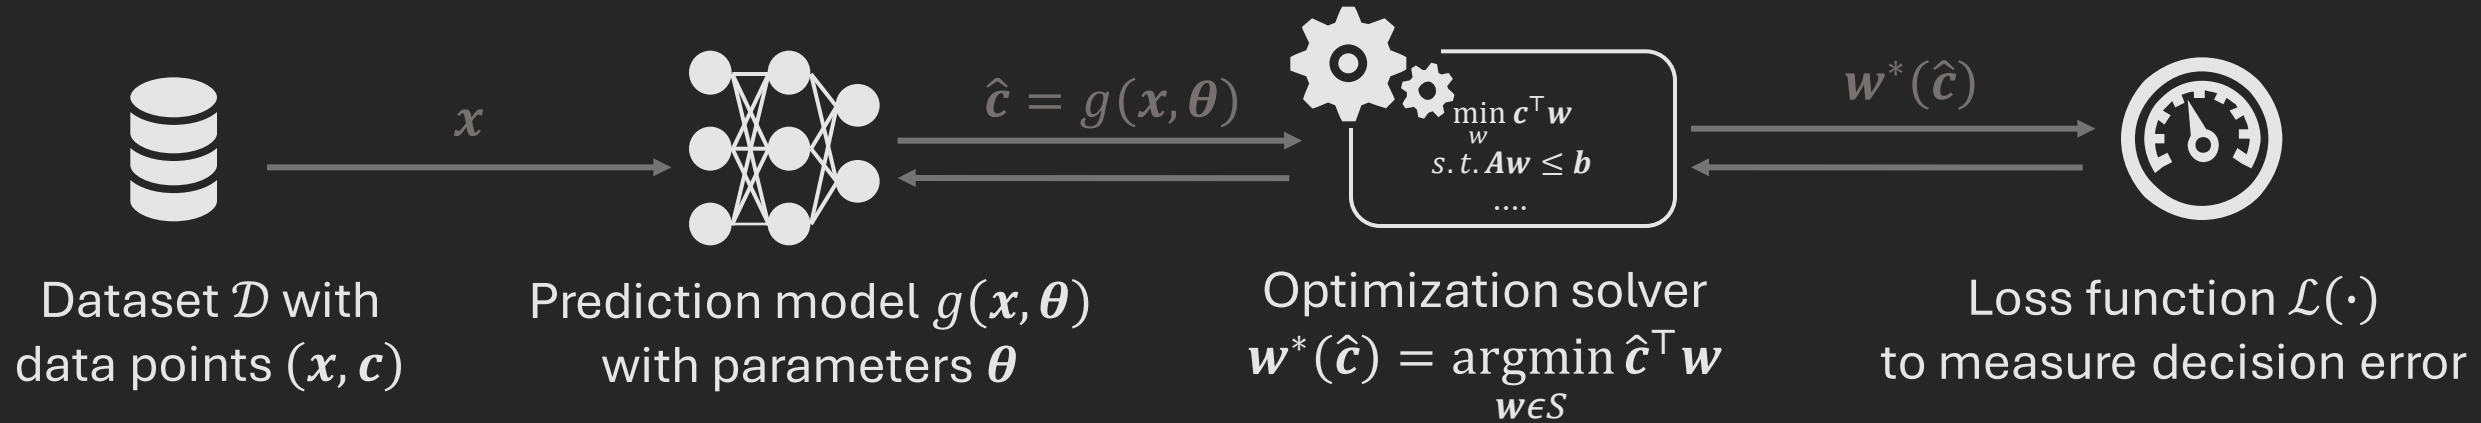
❖ Portfolio Optimization

? Unknown Costs: Travel Time, Electricity Prices, Asset Returns

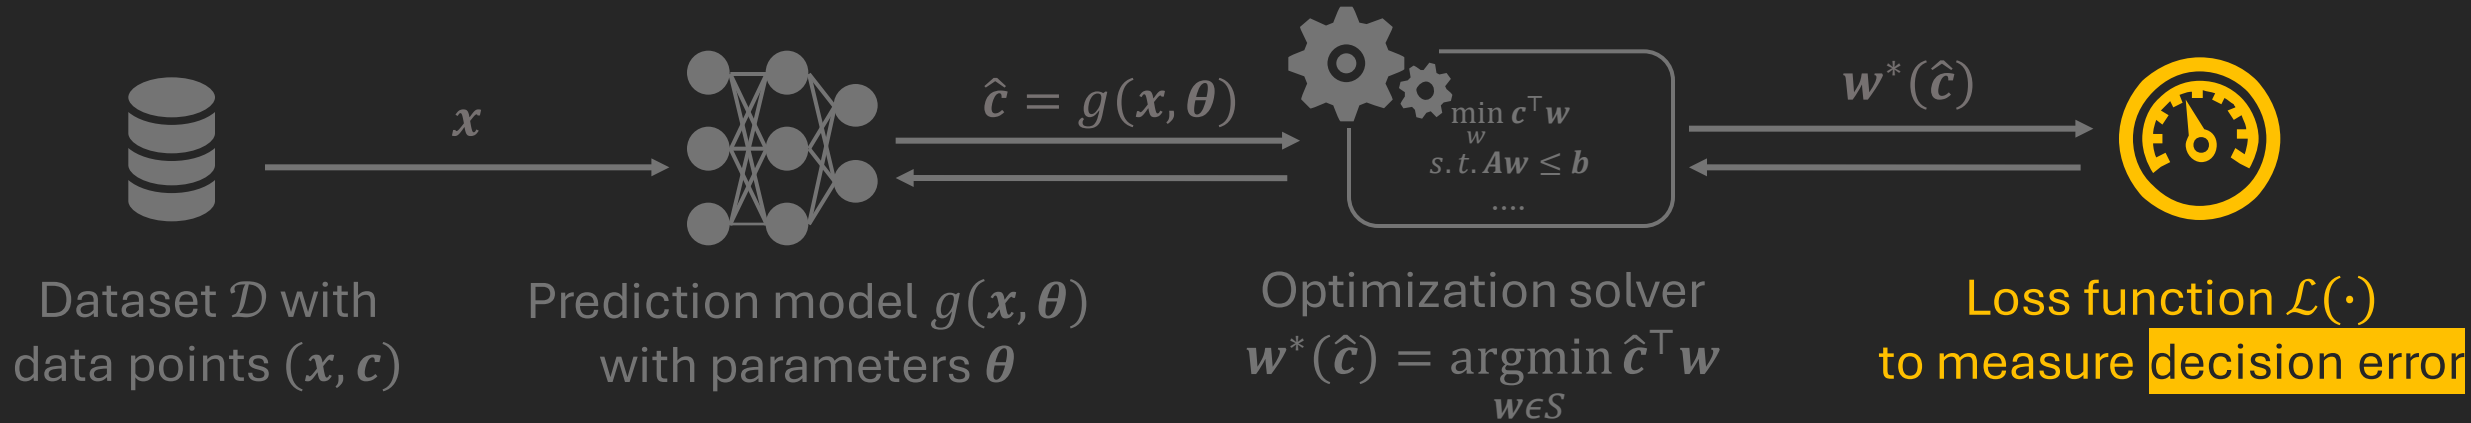


Observed Features: Distance, Time, Weather, Financial Factors...

Decision-focused Learning



Decision-focused Learning

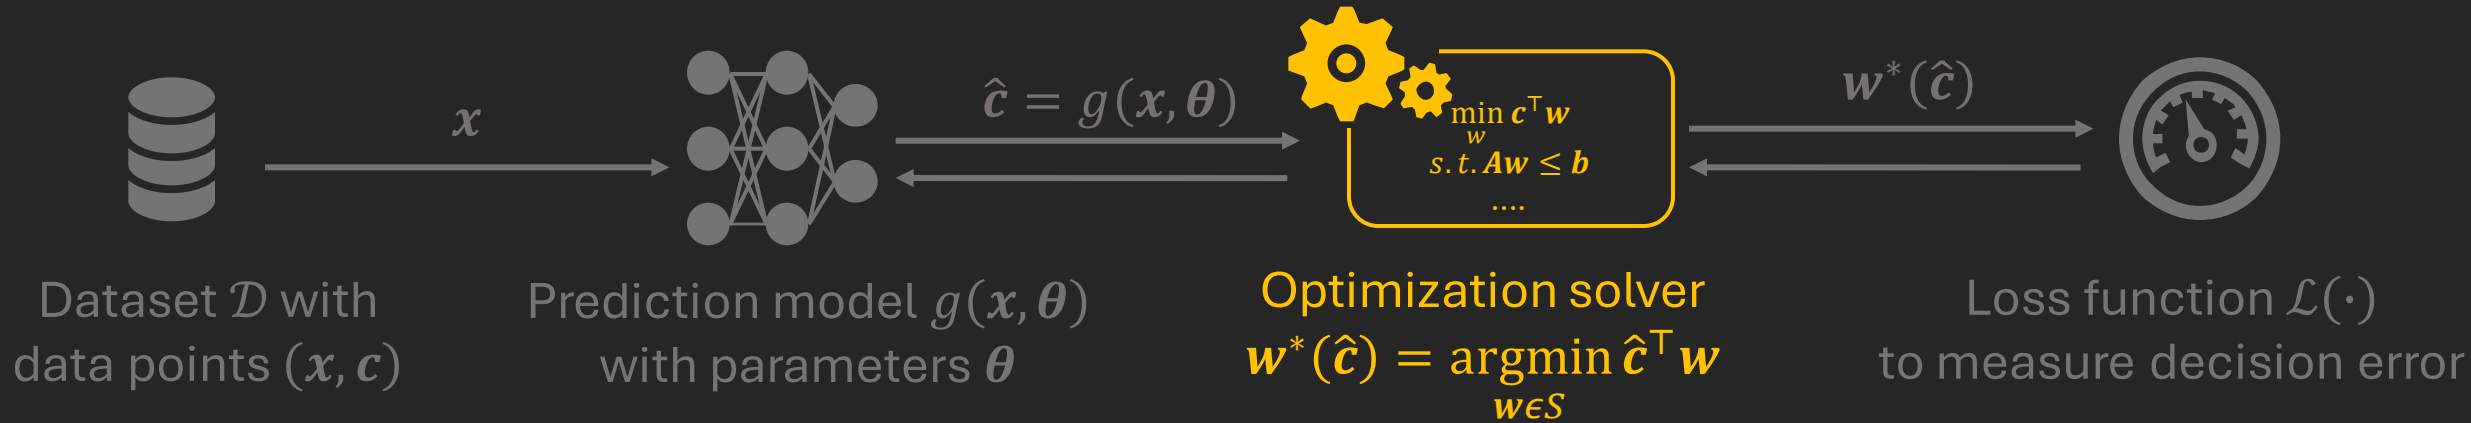


e.g.

$$\mathcal{L}_{\text{Regret}}(\hat{c}, c) = c^\top w^*(\hat{c}) - c^\top w^*(c)$$

$$\mathcal{L}_{\text{Square}}(\hat{c}, c) = \frac{1}{2} \|w^*(c) - w^*(\hat{c})\|_2^2$$

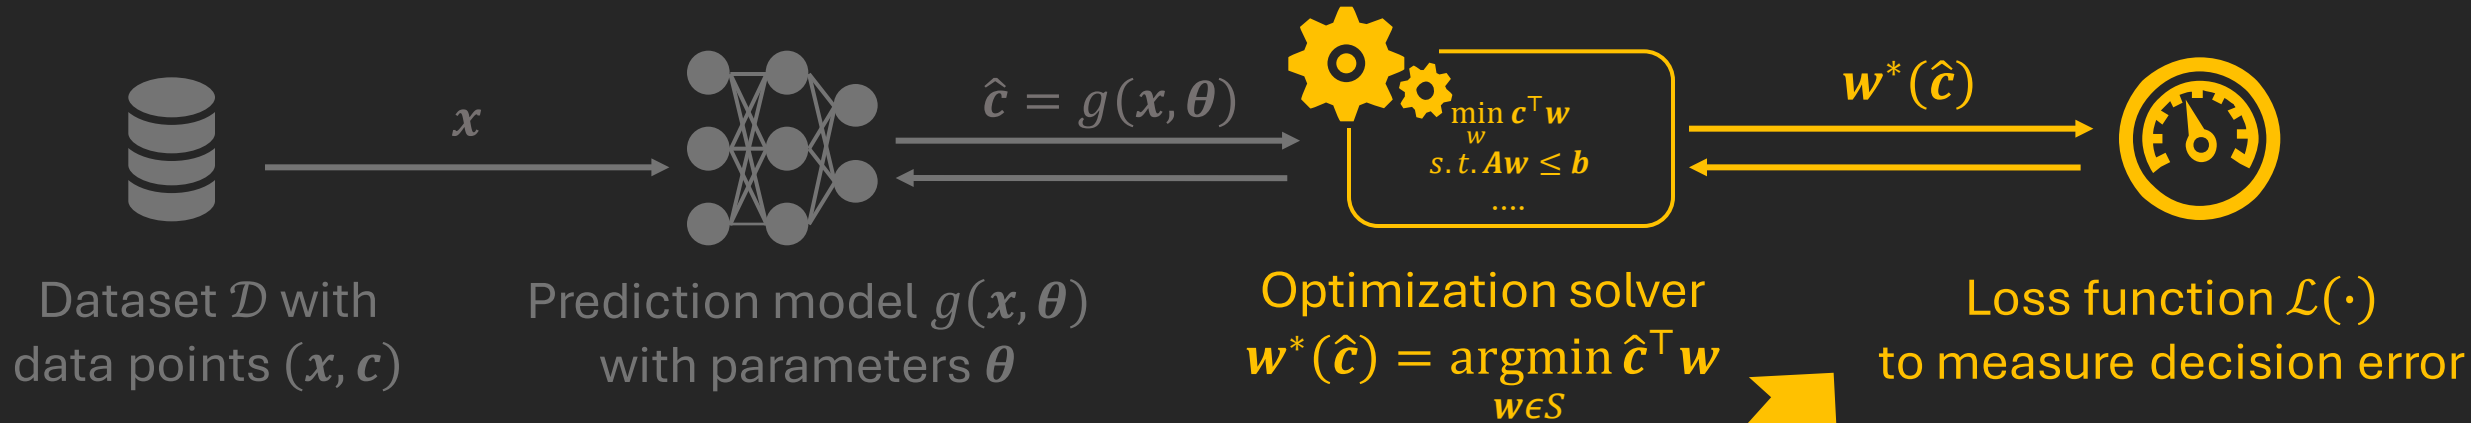
Decision-focused Learning



Computational Bottleneck:


All state-of-the-art methods require repeated solving during the iteration.

Decision-focused Learning



Cone-aligned Vector Estimation:
 Replace the original optimization problem
 with projection (quadratic programming).

Similarity Loss Function



When the predicted cost vector lies inside the **optimal subcone**, the optimal solution to the linear relaxation is the optimal of original BLP problem.

Similarity Loss Function



$$\mathcal{L}_{\text{CaVE}}(\hat{c}, c) = -\cos \phi$$

IMPORTANT QUESTIONS:

- What is the optimal subcone?
- How to obtain the angle ϕ ?

When the predicted cost vector lies inside the **optimal subcone**, the optimal solution to the linear relaxation is the optimal of original BLP problem.


Optimal Subcone



Binary Linear Program

1. For ILP, the normal cone to the convex hull at the optimal solution $w^*(c)$ is defined as **optimal cone** $\mathcal{C}^*(c)$.

Optimal Subcone



Binary Linear Program

1. For ILP, the normal cone to the convex hull at the optimal solution $w^*(c)$ is defined as **optimal cone** $\mathcal{C}^*(c)$.
2. $\forall c' \in \mathcal{C}^*(c), w^*(c') = w^*(c)$. Cost vectors yield the same optimal solution if and only if they are in the same optimal cone.


Optimal Subcone





Binary Linear Program

1. For ILP, the normal cone to the convex hull at the optimal solution $w^*(c)$ is defined as **optimal cone** $\mathcal{C}^*(c)$.
2. $\forall c' \in \mathcal{C}^*(c), w^*(c') = w^*(c)$. Cost vectors yield the same optimal solution if and only if they are in the same optimal cone.
3. However, for ILP, obtaining the convex hull is **NOT** trivial. e.g., Cutting Plane method...

Optimal Subcone




Binary Linear Program



Linear Relaxation

Optimal Subcone

1. For BLP, the normal cone to the feasible region of linear relaxation at the optimal solution $w^*(c)$ is defined as **optimal subcone** $\mathcal{SC}^*(c)$.




Linear Relaxation

Optimal Subcone

1. For BLP, the normal cone to the feasible region of linear relaxation at the optimal solution $w^*(c)$ is defined as **optimal subcone** $\mathcal{SC}^*(c)$.

Note: This does not apply to general ILP!




Linear Relaxation


Optimal Subcone

1. For BLP, the normal cone to the feasible region of linear relaxation at the optimal solution $w^*(c)$ is defined as **optimal subcone** $\mathcal{SC}^*(c)$.
2. $\mathcal{SC}^*(c) \subseteq \mathcal{C}^*(c)$. Thus, cost vectors yield the same optimal solution if they are in the same optimal subcone.


Optimal Cone
 $\mathcal{C}^*(c)$



Optimal Subcone
 $\mathcal{SC}^*(c)$




Optimal Subcone
 $\mathcal{SC}^*(c)$



Linear Relaxation

Optimal Subcone

1. For BLP, the normal cone to the feasible region of linear relaxation at the optimal solution $w^*(c)$ is defined as **optimal subcone** $\mathcal{SC}^*(c)$.
2. $\mathcal{SC}^*(c) \subseteq \mathcal{C}^*(c)$. Thus, cost vectors yield the same optimal solution if they are in the same optimal subcone.
3. The optimal subcone is the conic combination of tight constraints, which is trivial.

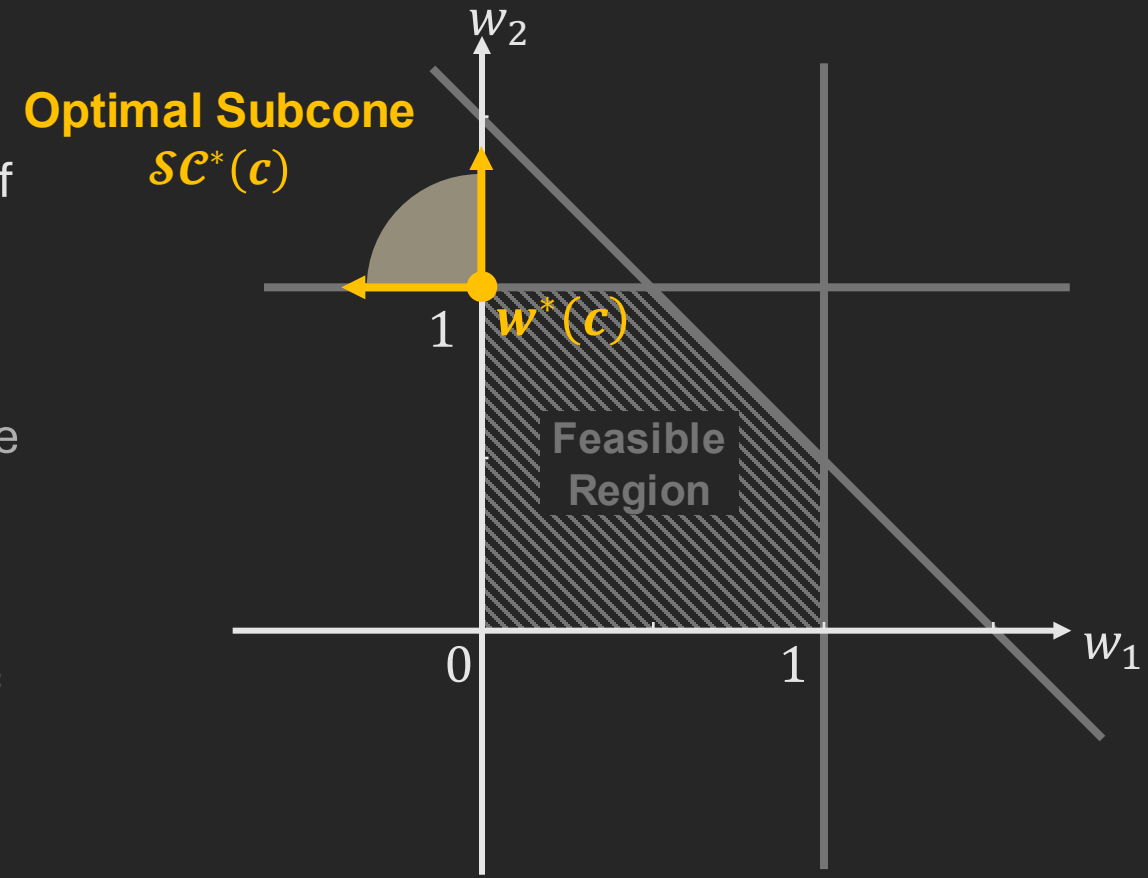


Linear Relaxation

Optimal Subcone

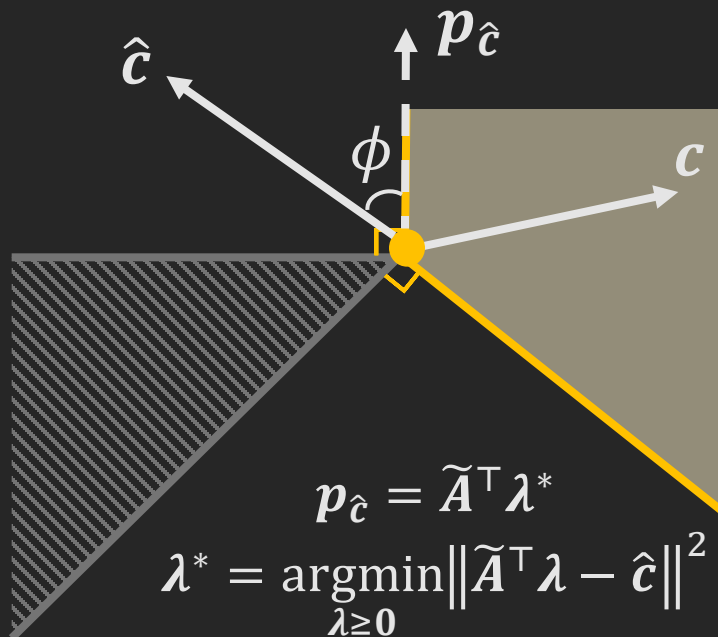
1. For BLP, the normal cone to the feasible region of linear relaxation at the optimal solution $w^*(c)$ is defined as **optimal subcone** $\mathcal{SC}^*(c)$.
2. $\mathcal{SC}^*(c) \subseteq \mathcal{C}^*(c)$. Thus, cost vectors yield the same optimal solution if they are in the same optimal subcone.
3. The optimal subcone is the conic combination of **tight constraints**, which is trivial.

$$\tilde{A}(c): \tilde{A}(c)^\top w^*(c) = b$$

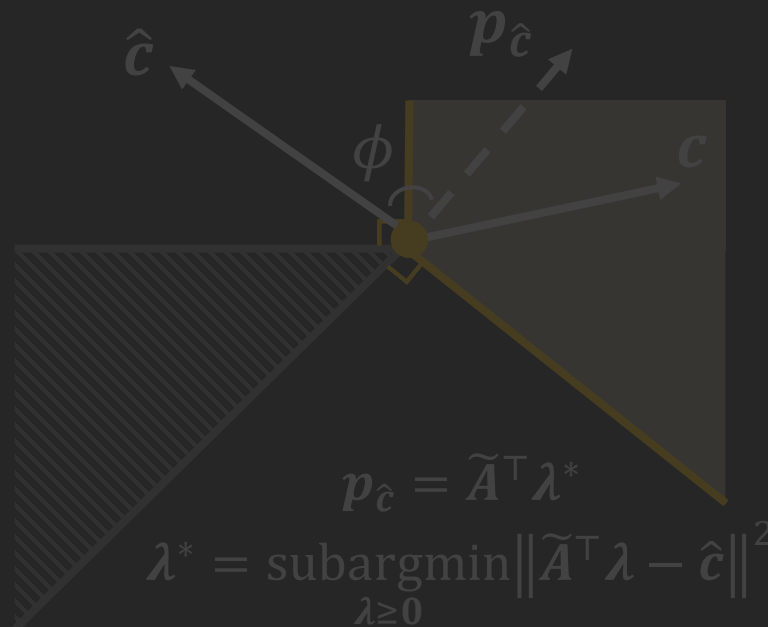


Linear Relaxation

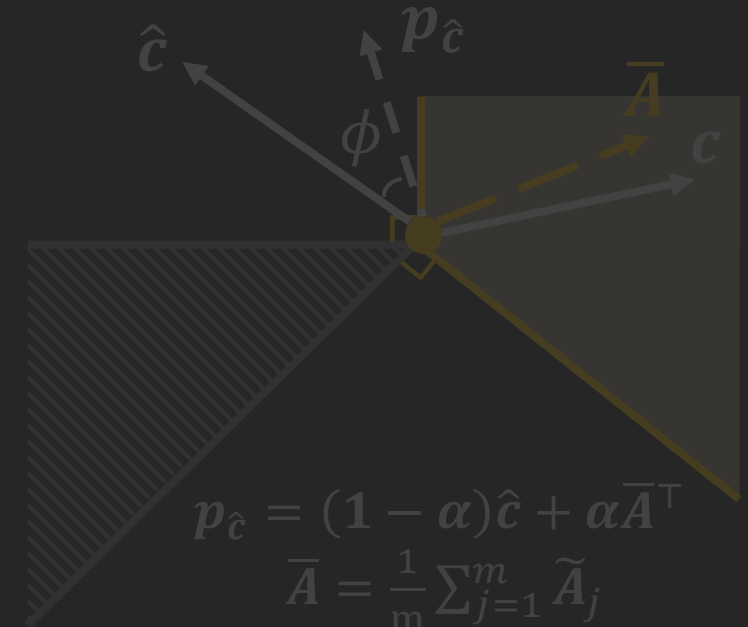
Projection



Exact Projection



Inner Projection




Heuristic Projection




CaVE Exact performs exact projection, wherein the optimal solution of the NNLS is on the surface of the cone.

This approach results in the vanishing gradients as the predicted cost vector close the surface.


Projection



Exact Projection



Inner Projection



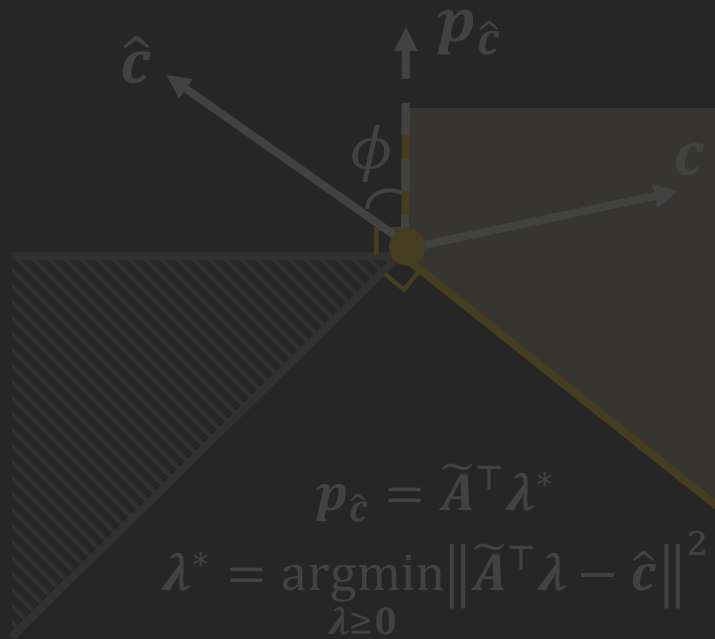
Heuristic Projection



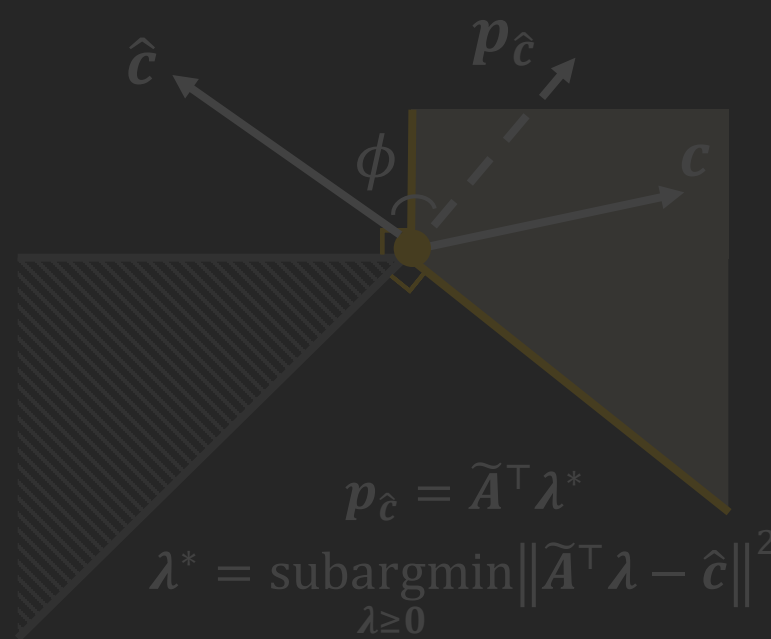
The goal is to obtain a projection of the predicted cost vector that lies inside the subcone. (suboptimal solution)

Since the solver (Clarabel) uses the primal-dual interior point, the feasibility is guaranteed at each iteration. We can simply set the maximum iterations.

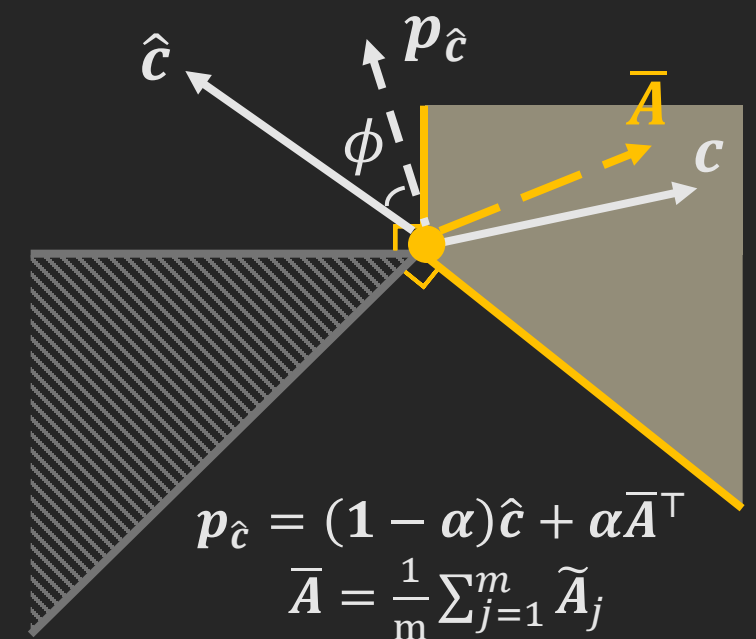
Projection



Exact Projection



Inner Projection



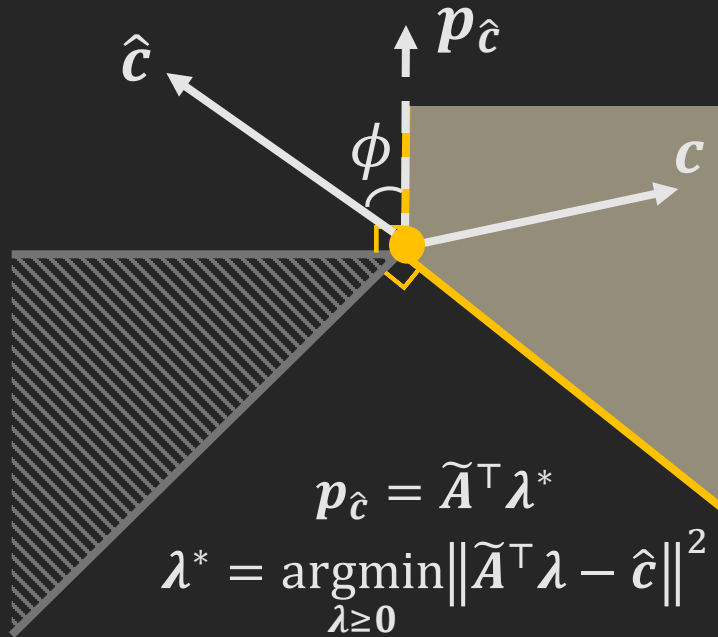
Heuristic Projection



A heuristic projection that does not require solving NNLS and relies on simple operations.

This approach does NOT guarantee feasibility, yet it ensures that the cost vector is pushed towards the cone.

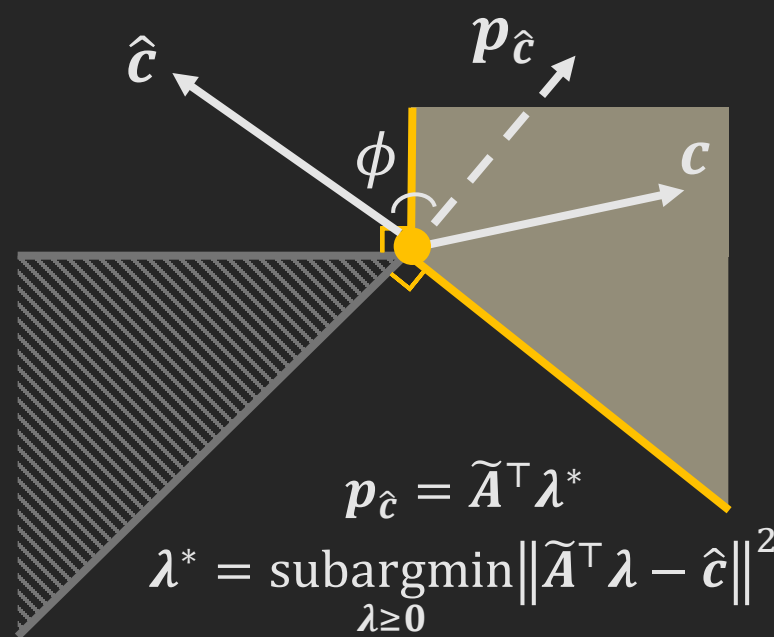
Projection



Exact Projection




CaVE Exact



Inner Projection



CaVE+



Heuristic Projection



CaVE Hybrid



Algorithm

Algorithm Cone-aligned Vector Estimation (CaVE)

Require: Pairs of feature vectors and binding constraints $\{(\mathbf{x}^i, \tilde{\mathbf{A}}^i)\}_{i=1}^n$ for n training instances; learning rate $\alpha > 0$

```
1: Initialize model parameters  $\theta$ 
2: for each training epoch do
3:   for each batch of training samples  $(\mathbf{x}, \tilde{\mathbf{A}})$  do
4:     Predict cost coefficient  $\hat{\mathbf{c}} \leftarrow g(\mathbf{x}, \theta)$ 
5:     Compute projection  $\mathbf{p}_{\hat{\mathbf{c}}}$  with quadratic program
6:     Compute cosine similarity loss  $\mathcal{L}_{\text{CaVE}}(\hat{\mathbf{c}}, \tilde{\mathbf{A}})$ 
7:     Compute the gradient  $\nabla_{\theta} \mathcal{L}_{\text{CaVE}}(\hat{\mathbf{c}}, \tilde{\mathbf{A}})$  with backpropagation
8:     Update ML model parameters  $\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{CaVE}}(\hat{\mathbf{c}}, \tilde{\mathbf{A}})$ 
9:   end for
10: end for
11: return  $g(\cdot, \theta)$ 
```

Algorithm

Algorithm Cone-aligned Vector Estimation (CaVE)

Require: Pairs of feature vectors and binding constraints $\{(\mathbf{x}^i, \tilde{\mathbf{A}}^i)\}_{i=1}^n$ for n training instances; learning rate $\alpha > 0$

```
1: Initialize model parameters  $\theta$ 
2: for each training epoch do
3:   for each batch of training samples  $(\mathbf{x}, \tilde{\mathbf{A}})$  do
4:     Predict cost coefficient  $\hat{\mathbf{c}} \leftarrow g(\mathbf{x}, \theta)$ 
5:     Compute projection  $\mathbf{p}_{\hat{\mathbf{c}}}$  with quadratic program
6:     Compute cosine similarity loss  $\mathcal{L}_{\text{CaVE}}(\hat{\mathbf{c}}, \tilde{\mathbf{A}})$ 
7:     Compute the gradient  $\nabla_{\theta} \mathcal{L}_{\text{CaVE}}(\hat{\mathbf{c}}, \tilde{\mathbf{A}})$  with backpropagation
8:     Update ML model parameters  $\theta \leftarrow \theta - \alpha \nabla_{\theta} \mathcal{L}_{\text{CaVE}}(\hat{\mathbf{c}}, \tilde{\mathbf{A}})$ 
9:   end for
10: end for
11: return  $g(\cdot, \theta)$ 
```

Faster than solving BLP!

Experiments - SP5

Shortest Path on 5 × 5 Grid

Average Test **Normalized Regret** (%) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	8.82 ± 1.15	10.73 ± 1.54	8.39 ± 0.95	8.35 ± 0.88	7.79 ± 1.00	7.68 ± 0.99	11.34 ± 1.11
Deg 6	12.58 ± 2.14	11.30 ± 1.30	8.89 ± 0.90	8.84 ± 1.00	7.72 ± 1.11	7.86 ± 0.96	13.78 ± 1.58

Average **Training Time** (Sec) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	1.52 ± 0.14	4.64 ± 0.09	4.89 ± 0.12	2.57 ± 0.19	17.64 ± 0.12	18.52 ± 0.31	4.50 ± 0.48
Deg 6	1.38 ± 0.13	3.52 ± 0.11	3.72 ± 0.14	2.39 ± 0.19	18.68 ± 0.40	17.78 ± 0.13	4.38 ± 0.42

Experiments – TSP20

Traveling Salesperson with 20 Nodes

Average Test **Normalized Regret** (%) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	12.12 ± 0.89	7.35 ± 0.40	6.20 ± 0.24	7.69 ± 0.33	5.95 ± 0.16	6.56 ± 0.21	12.21 ± 0.88
Deg 6	21.32 ± 1.81	8.01 ± 0.45	6.97 ± 0.37	9.52 ± 0.64	7.48 ± 0.36	7.41 ± 0.37	14.31 ± 0.40

Average **Training Time** (Sec) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	1.52 ± 0.10	113.56 ± 3.16	107.15 ± 3.80	27.06 ± 2.17	175.23 ± 4.95	220.21 ± 24.20	25.92 ± 4.23
Deg 6	1.53 ± 0.19	158.66 ± 9.65	102.19 ± 10.38	30.17 ± 2.62	185.13 ± 7.44	185.02 ± 5.09	25.48 ± 3.66

Experiments – TSP50

Traveling Salesperson with 50 Nodes

Average Test **Normalized Regret** (%) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	28.16 ± 1.08	15.19 ± 0.65	7.69 ± 0.22	9.59 ± 0.44	7.57 ± 0.20	8.03 ± 0.23	14.31 ± 0.40
Deg 6	52.61 ± 2.36	23.25 ± 2.41	8.57 ± 0.38	11.28 ± 0.80	10.26 ± 0.46	9.00 ± 0.52	17.12 ± 0.48

Average **Training Time** (Sec) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	1.55 ± 0.18	611.47 ± 23.52	518.07 ± 51.89	196.96 ± 35.92	1220.68 ± 85.39	1328.99 ± 28.87	151.80 ± 24.21
Deg 6	1.16 ± 0.13	502.71 ± 16.03	573.87 ± 20.19	253.93 ± 27.67	1191.29 ± 42.63	1456.21 ± 34.18	155.95 ± 24.46

Experiments – CVRP20

Capacity Vehicle Routing with 20 Nodes

Average Test **Normalized Regret** (%) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	10.10 ± 0.64	9.26 ± 1.56	6.44 ± 0.24	7.92 ± 0.52	5.94 ± 0.25	6.32 ± 0.28	15.77 ± 0.96
Deg 6	19.50 ± 1.22	11.64 ± 0.25	7.94 ± 0.54	11.44 ± 1.14	8.75 ± 0.28	8.09 ± 0.57	18.96 ± 1.01

Average **Training Time** (Sec) with Standard Deviation

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	1.65 ± 0.48	213.56 ± 42.36	153.56 ± 11.08	44.52 ± 6.27	7020.11 ± 1043.05	3773.31 ± 288.84	583.56 ± 170.67
Deg 6	1.54 ± 0.25	208.95 ± 12.90	127.94 ± 13.84	51.83 ± 8.78	2204.83 ± 99.86	6197.84 ± 288.63	470.20 ± 84.46

Experiments – CVRP30

Capacity Vehicle Routing with 30 Nodes

Test **Normalized Regret** (%)

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	19.72	12.54	9.13	9.99	N/A		18.28

Training Time (Sec)

Methods	2-Stage	CaVE-E	CaVE+	CaVE-H	SPO+	PFYL	NCE
Deg 4	9.27	331.73	287.77	132.62	≥100h		884.95

Due to the scale of the problem, we did not repeat our experimental evaluation with random seeds.

Experiments – Relaxation

Traveling Salesperson with 50 Nodes

Average Test **Normalized Regret** (%)

Methods	CaVE+	SPO+ Rel	PFYL Rel
Deg 4	7.69 \pm 0.22	10.17 \pm 0.23	11.11 \pm 0.33
Deg 6	8.57 \pm 0.38	13.14 \pm 0.46	13.38 \pm 0.58

Average **Training Time** (Sec)

Methods	CaVE+	SPO+ Rel	PFYL Rel
Deg 4	518.07 \pm 51.89	386.06 \pm 9.69	536.67 \pm 4.94
Deg 6	573.87 \pm 20.19	636.99 \pm 3.04	510.37 \pm 3.46

Capacity Vehicle Routing with 20 Nodes

Average Test **Normalized Regret** (%)

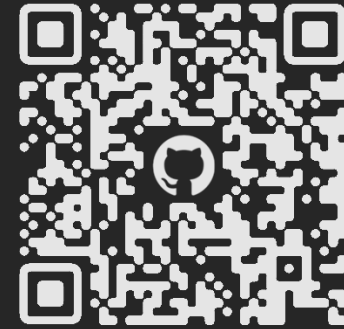
Methods	CaVE+	SPO+ Rel	PFYL Rel
Deg 4	6.44 \pm 0.24	8.03 \pm 0.38	17.07 \pm 0.63
Deg 6	7.94 \pm 0.54	15.73 \pm 0.39	19.19 \pm 1.66

Average **Training Time** (Sec)

Methods	CaVE+	SPO+ Rel	PFYL Rel
Deg 4	153.56 \pm 11.08	78.95 \pm 0.73	78.80 \pm 1.19
Deg 6	127.94 \pm 13.84	78.74 \pm 3.82	81.80 \pm 0.86

Thank You

GitHub



Mechanical & Industrial Engineering
UNIVERSITY OF TORONTO